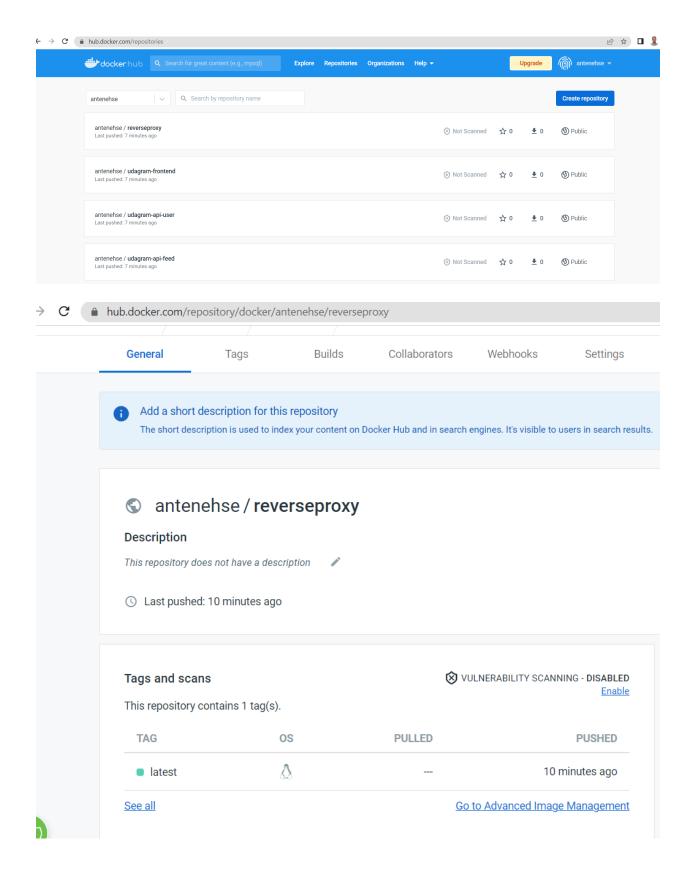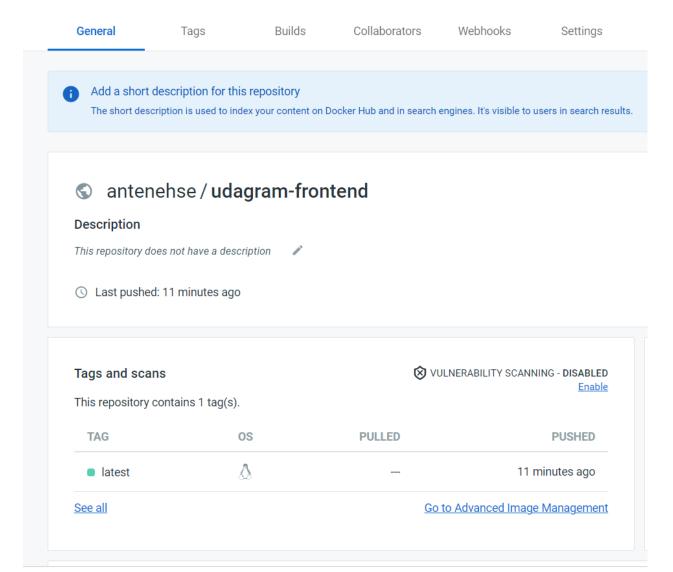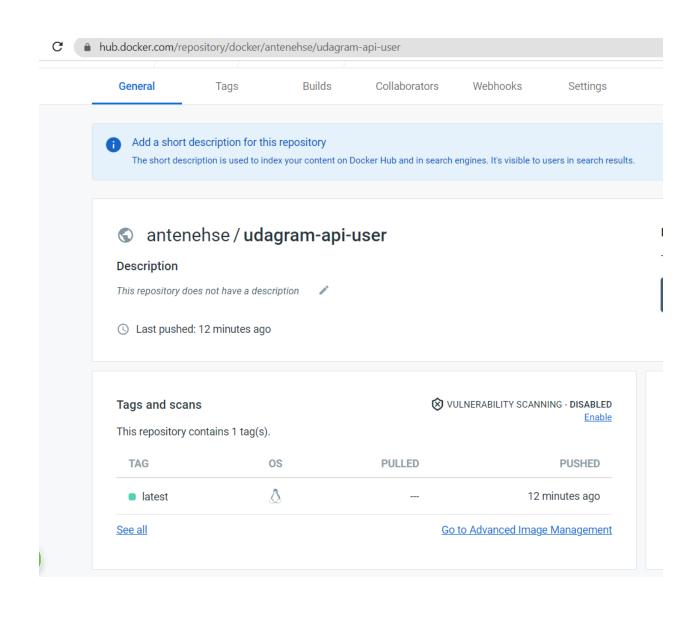- Docker images in your repository in DockerHub

antenehse ▾   Search by repository name   Create repository

antenehse / **reverseproxy**
Last pushed: 7 minutes ago
⊗ Not Scanned   ☆ 0   ⬇ 0   Ⓢ Public

antenehse / **udagram-frontend**
Last pushed: 7 minutes ago
⊗ Not Scanned   ☆ 0   ⬇ 0   Ⓢ Public

antenehse / **udagram-api-user**
Last pushed: 7 minutes ago
⊗ Not Scanned   ☆ 0   ⬇ 0   Ⓢ Public

antenehse / **udagram-api-feed**
Last pushed: 7 minutes ago
⊗ Not Scanned   ☆ 0   ⬇ 0   Ⓢ Public

General   Tags   Builds   Collaborators   Webhooks   Settings

ℹ️ **Add a short description for this repository**
The short description is used to index your content on Docker Hub and in search engines. It's visible to users in search results.

🌐 **antenehse / reverseproxy**

Description

This repository does not have a description   ✏️

🕐 Last pushed: 10 minutes ago

Tags and scans   ⊗ VULNERABILITY SCANNING - **DISABLED**
                                                                                    Enable
This repository contains 1 tag(s).

TAG	OS	PULLED	PUSHED
● latest	🐧	---	10 minutes ago

See all   Go to Advanced Image Management

ℹ **Add a short description for this repository**
The short description is used to index your content on Docker Hub and in search engines. It's visible to users in search results.

🌐 **antenehse / udagram-frontend**

Description

This repository does not have a description   ✏

🕐 Last pushed: 11 minutes ago

Tags and scans   🛡 VULNERABILITY SCANNING - **DISABLED**
                                                                              Enable
This repository contains 1 tag(s).

TAG	OS	PULLED	PUSHED
🟢 latest	🐧	---	11 minutes ago

See all   Go to Advanced Image Management

General   Tags   Builds   Collaborators   Webhooks   Settings

ℹ **Add a short description for this repository**
The short description is used to index your content on Docker Hub and in search engines. It's visible to users in search results.

🌐 **antenehse / udagram-api-user**

Description

This repository does not have a description ✏

🕐 Last pushed: 12 minutes ago

Tags and scans   🛡 VULNERABILITY SCANNING - **DISABLED**
                                                                                Enable
This repository contains 1 tag(s).

TAG	OS	PULLED	PUSHED
● latest	🐧	---	12 minutes ago

See all   Go to Advanced Image Management

TravisCI build pipeline showing successful build jobs

Screenshot 1 (top):

app.circleci.com/pipelines/github/Anteneh2121/Microservice?branch=main

Anteneh2121
Anteneh Kassaw Bizu...

Dashboard
Projects
Insights
Organization Settings
Plan

CI behind your firewall just got easier

A more scalable, container friendly self-hosted runner is now in open preview

⚠ CircleCI Security Alert: There is an ongoing phishing attempt to obtain CircleCI credentials. For more information, please see our post here ✕

Dashboard | Project | Branch
All Pipelines > Microservice > main

Microservice ⚏ Add team members

Edit Config | Trigger Pipeline | Project Settings

Filters
Everyone's Pipelines ▾ | Microservice ▾ | main ▾ | All days ▾ | ▾

Auto-expand ⬤

Pipeline	Status	Workflow	Branch / Commit	Start	Duration	Actions
Microservice 1 ▸	✔ Success	default	main aa5c61b	17m ago	3m 5s	↻ ↻ ✕ ⋯
▸	⚠ Failed	default	main aa5c61b	30m ago	3m 50s	↻ ↻ ✕ ⋯

Screenshot 2 (bottom):

app.circleci.com/pipelines/github/Anteneh2121/Microservice?branch=main

Anteneh2121
Anteneh Kassaw Bizu...

Dashboard
Projects
Insights
Organization Settings
Plan

CI behind your firewall just got easier

A more scalable, container friendly self-hosted runner is now in open preview

⚠ CircleCI Security Alert: There is an ongoing phishing attempt to obtain CircleCI credentials. For more information, please see our post here ✕

Dashboard | Project | Branch
All Pipelines > Microservice > main

Microservice ⚏ Add team members

Edit Config | Trigger Pipeline | Project Settings

Filters
Everyone's Pipelines ▾ | Microservice ▾ | main ▾ | All days ▾ | ▾

Auto-expand ⬤

Pipeline	Status	Workflow	Branch / Commit	Start	Duration	Actions
Microservice 1 ▾	✔ Success	default	main aa5c61b	18m ago	3m 5s	↻ ↻ ✕ ⋯
Jobs	✔ lint-app 3				3s	
	✔ build-docker 4				2m 57s	

Anteneh2121
Anteneh Kassaw Bizu...

Dashboard   Project   Branch   Workflow   Job

All Pipelines > Microservice > main > default > build-docker (4)

Dashboard

Projects

Insights

Organization Settings

Plan

build-docker ✓ Success

Duration / Finished	Queued	Executor / Resource Class	Branch	Commit	Author
🕐 2m 56s / 15m ago	⧗ 0s	🐧 Machine / Linux Large ↗ ⓘ	⅛ main		👤 Anteneh Kassaw Bizuneh

CI behind your firewall just got easier

A more scalable, **container friendly self-hosted runner** is now in open preview

Getting Started

Notifications

Status OPERATIONAL

STEPS   TESTS   TIMING   ARTIFACTS   RESOURCES NEW

Parallel runs

📦	0 02:56	📦	Use parallelism to run faster tests	Go to Docs	✕
			Parallelism speeds up tests by splitting them across multiple executors.		

▶ ✓ Spin up environment

▶ ✓ Preparing environment variables

▶ ✓ Checkout code

▶ ✓ Build docker container for each microservices

Anteneh2121
Anteneh Kassaw Bizu...

- Dashboard
- Projects
- Insights
- Organization Settings
- Plan

CI behind your firewall just got easier

A more scalable, container friendly self-hosted runner is now in open preview

- Getting Started
- Notifications
- Status `OPERATIONAL`

```
3674
3675   Login Succeeded
3676   -------------- Pushing images to dockerhub-------
3677   Using ******* tag: latest
3678   The push refers to repository [docker.io/*********/udagram-api-feed]
3679
3680   e521c27a1f13: Preparing
3681   40c539836a8d: Preparing
3682   b5a945995f91: Preparing
3683   a617ef186396: Preparing
3684   586c0b414da7: Preparing
3685   0bfd290f2c17: Preparing
3686   6d75cd01c26c: Preparing
3687   95904c181913: Preparing
3688   df69bfa94785: Preparing
3689   f35deb8d96fc: Preparing
3690   f6c2459e2059: Preparing
3691   f8323fb3a55c: Preparing
3692   latest: digest: sha256:021f02f76b717df1b1662720a60bb770c3a4c1c3bfee954b3765657c2ff703a8 size: 3054
3693   Using ******* tag: latest
3694   The push refers to repository [docker.io/*********/udagram-api-user]
3695
3696   a20ce594614d: Preparing
3697   40c539836a8d: Preparing
3698   b5a945995f91: Preparing
3699   a617ef186396: Preparing
3700   586c0b414da7: Preparing
3701   0bfd290f2c17: Preparing
3702   6d75cd01c26c: Preparing
3703   95904c181913: Preparing
3704   df69bfa94785: Preparing
3705   f35deb8d96fc: Preparing
```

```
3698   b5a945995f91: Preparing
3699   a617ef186396: Preparing
3700   586c0b414da7: Preparing
3701   0bfd290f2c17: Preparing
3702   6d75cd01c26c: Preparing
3703   95904c181913: Preparing
3704   df69bfa94785: Preparing
3705   f35deb8d96fc: Preparing
3706   f6c2459e2059: Preparing
3707   f8323fb3a55c: Preparing
3708   latest: digest: sha256:2c3d5acb246a825c6d5b024f8e841105705a0559d97ab5f1b444f2a318b9961d size: 3054
3709   Using ******* tag: latest
3710   The push refers to repository [docker.io/*********/udagram-frontend]
3711
3712   af1d212957ec: Preparing
3713   bf4e176a4d9b: Preparing
3714   a1d571e4e83d: Preparing
3715   6d97b4d00719: Preparing
3716   2a7647ca3937: Preparing
3717   549c42eea4a6: Preparing
3718   latest: digest: sha256:23cccfd824700ad972f5095bd5edccebb785194aaa1be799184ee5b6d28f04b6 size: 1779
3719   Using ******* tag: latest
3720   The push refers to repository [docker.io/*********/reverseproxy]
3721
3722   0ced324d1f1d: Preparing
3723   bf4e176a4d9b: Preparing
3724   a1d571e4e83d: Preparing
3725   6d97b4d00719: Preparing
3726   2a7647ca3937: Preparing
3727   549c42eea4a6: Preparing
3728   latest: digest: sha256:da06b4e9f0de737998d397b4dddfba0d934fb11af64b0db4eeeeba2ed5a7fa6e size: 1775
3729   CircleCI received exit code 0
```

Kubernetes `kubectl get pods` output

```
hp@ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$ kubectl get pods
NAME                                READY   STATUS    RESTARTS        AGE
backend-feed-7bc58f6dcb-csbd9       1/1     Running   0               6s
backend-feed-7bc58f6dcb-fktd2       1/1     Running   0               3s
backend-user-67c44cd647-8sj4q       1/1     Running   1 (15s ago)     75s
backend-user-67c44cd647-gkjfc       1/1     Running   1 (13s ago)     73s
frontend-75db855b87-8v8cb           1/1     Running   0               11m
frontend-75db855b87-m6252           1/1     Running   0               11m
reverseproxy-7c8bbd75c4-ppxg7       1/1     Running   3 (21m ago)     22m
reverseproxy-7c8bbd75c4-x97t4       1/1     Running   3 (21m ago)     22m

hp@ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$ |
```

```
hp@ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$ kubectl get pods
NAME                                READY   STATUS    RESTARTS        AGE
backend-feed-76c9858588-gztg6       1/1     Running   0               26s
backend-feed-76c9858588-vtr5q       1/1     Running   0               24s
backend-user-7cbd499f75-f2jst       1/1     Running   0               11s
backend-user-7cbd499f75-rwf9p       1/1     Running   0               9s
frontend-65554cd4f5-n45rj           1/1     Running   0               25h
frontend-65554cd4f5-pq7jf           1/1     Running   0               25h
reverseproxy-7c8bbd75c4-ppxg7       1/1     Running   3 (26h ago)     26h
reverseproxy-7c8bbd75c4-x97t4       1/1     Running   3 (26h ago)     26h
```

Kubernetes `kubectl describe services` output

```
hp@ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$ kubectl describe services
Name:              backend-feed
Namespace:         default
Labels:            service=backend-feed
Annotations:       <none>
Selector:          service=backend-feed
Type:              ClusterIP
IP Family Policy:  SingleStack
IP Families:       IPv4
IP:                10.100.227.230
IPs:               10.100.227.230
Port:              8080   8080/TCP
TargetPort:        8080/TCP
Endpoints:         <none>
Session Affinity:  None
Events:            <none>


Name:              backend-user
Namespace:         default
Labels:            service=backend-user
Annotations:       <none>
Selector:          service=backend-user
Type:              ClusterIP
IP Family Policy:  SingleStack
IP Families:       IPv4
IP:                10.100.99.118
IPs:               10.100.99.118
Port:              8080   8080/TCP
TargetPort:        8080/TCP
Endpoints:         192.168.33.7:8080,192.168.38.27:8080,192.168.4.193:8080
Session Affinity:  None
Events:            <none>


Name:              frontend
Namespace:         default
Labels:            service=frontend
Annotations:       <none>
Selector:          service=frontend
Type:              ClusterIP
IP Family Policy:  SingleStack
IP Families:       IPv4
IP:                10.100.151.69
IPs:               10.100.151.69
Port:              8100   8100/TCP
TargetPort:        80/TCP
Endpoints:         192.168.1.37:80,192.168.11.243:80,192.168.35.152:80
Session Affinity:  None
Events:            <none>


Name:              kubernetes
Namespace:         default
Labels:            component=apiserver
                   provider=kubernetes
```

```
Endpoints:              192.168.1.37:80,192.168.11.243:80,192.168.35.152:80
Session Affinity:       None
Events:                 <none>


Name:                   kubernetes
Namespace:              default
Labels:                 component=apiserver
                        provider=kubernetes
Annotations:            <none>
Selector:               <none>
Type:                   ClusterIP
IP Family Policy:       SingleStack
IP Families:            IPv4
IP:                     10.100.0.1
IPs:                    10.100.0.1
Port:                   https   443/TCP
TargetPort:             443/TCP
Endpoints:              192.168.102.115:443,192.168.128.166:443
Session Affinity:       None
Events:                 <none>


Name:                   publicfrontend
Namespace:              default
Labels:                 service=frontend
Annotations:            <none>
Selector:               app=frontend
Type:                   LoadBalancer
IP Family Policy:       SingleStack
IP Families:            IPv4
IP:                     10.100.41.227
IPs:                    10.100.41.227
LoadBalancer Ingress:   a7c725c6ba4514a1bb67987bb7fa47f9-599922655.us-east-2.elb.amazonaws.com
Port:                   <unset>  80/TCP
TargetPort:             80/TCP
NodePort:               <unset>  30896/TCP
Endpoints:              192.168.1.37:80,192.168.11.243:80,192.168.35.152:80
Session Affinity:       None
External Traffic Policy: Cluster
Events:                 <none>


Name:                   publicreverseproxy
Namespace:              default
Labels:                 service=reverseproxy
Annotations:            <none>
Selector:               app=reverseproxy
Type:                   LoadBalancer
IP Family Policy:       SingleStack
IP Families:            IPv4
IP:                     10.100.252.191
IPs:                    10.100.252.191
LoadBalancer Ingress:   a1cce38d4be444e66a186c06df52dffc-662956085.us-east-2.elb.amazonaws.com
Port:                   <unset>  8080/TCP
TargetPort:             8080/TCP
NodePort:               <unset>  31839/TCP
Endpoints:              192.168.25.89:8080,192.168.33.42:8080,192.168.55.204:8080
```

```
Name:                   publicreverseproxy
Namespace:              default
Labels:                 service=reverseproxy
Annotations:            <none>
Selector:               app=reverseproxy
Type:                   LoadBalancer
IP Family Policy:       SingleStack
IP Families:            IPv4
IP:                     10.100.252.191
IPs:                    10.100.252.191
LoadBalancer Ingress:   a1cce38d4be444e66a186c06df52dffc-662956085.us-east-2.elb.amazonaws.com
Port:                   <unset>  8080/TCP
TargetPort:             8080/TCP
NodePort:               <unset>  31839/TCP
Endpoints:              192.168.25.89:8080,192.168.33.42:8080,192.168.55.204:8080
Session Affinity:       None
External Traffic Policy: Cluster
Events:                 <none>


Name:                   reverseproxy
Namespace:              default
Labels:                 service=reverseproxy
Annotations:            <none>
Selector:               service=reverseproxy
Type:                   ClusterIP
IP Family Policy:       SingleStack
IP Families:            IPv4
IP:                     10.100.229.158
IPs:                    10.100.229.158
Port:                   8080  8080/TCP
TargetPort:             8080/TCP
Endpoints:              192.168.25.89:8080,192.168.33.42:8080,192.168.55.204:8080
Session Affinity:       None
Events:                 <none>

hp@ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$ |
```

Kubernetes `kubectl describe hpa` output

Kubectl get hpa

```
hp@ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$ kubectl get hpa
NAME            REFERENCE                TARGETS         MINPODS  MAXPODS  REPLICAS  AGE
backend-feed    Deployment/backend-feed  <unknown>/70%   3        5        3         39m
backend-user    Deployment/backend-user  <unknown>/70%   3        5        3         39m
frontend        Deployment/frontend      <unknown>/70%   3        5        3         39m
reverseproxy    Deployment/reverseproxy  <unknown>/70%   3        5        3         38m

hp@ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$ |
```

kubectl describe hpa

```
hp9ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$ kubectl describe hpa
Name:                                                    backend-feed
Namespace:                                               default
Labels:                                                  <none>
Annotations:                                             <none>
CreationTimestamp:                                       Sat, 24 Sep 2022 12:02:14 -0700
Reference:                                               Deployment/backend-feed
Metrics:                                                 ( current / target )
  resource cpu on pods  (as a percentage of request):   <unknown> / 70%
Min replicas:                                            3
Max replicas:                                            5
Deployment pods:                                         3 current / 3 desired
Conditions:
  Type           Status  Reason              Message
  ----           ------  ------              -------
  AbleToScale    True    SucceededGetScale   the HPA controller was able to get the target's current scale
  ScalingActive  False   FailedGetResourceMetric   the HPA was unable to compute the replica count: failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch metrics from resource me
trics API: the server could not find the requested resource (get pods.metrics.k8s.io)
Events:
  Type     Reason                   Age           From                       Message
  ----     ------                   ----          ----                       -------
  Normal   SuccessfulRescale        73s           horizontal-pod-autoscaler  New size: 3; reason: Current number of replicas below Spec.MinReplicas
  Warning  FailedGetResourceMetric  13s (x4 over 58s)  horizontal-pod-autoscaler  failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch metrics from resource metrics API: t
he server could not find the requested resource (get pods.metrics.k8s.io)
  Warning  FailedComputeMetricsReplicas  13s (x4 over 58s)  horizontal-pod-autoscaler  invalid metrics (1 invalid out of 1), first error is: failed to get cpu utilization: unable to get metrics for resource cpu
: unable to fetch metrics from resource metrics API: the server could not find the requested resource (get pods.metrics.k8s.io)


Name:                                                    backend-user
Namespace:                                               default
Labels:                                                  <none>
Annotations:                                             <none>
CreationTimestamp:                                       Sat, 24 Sep 2022 12:02:32 -0700
Reference:                                               Deployment/backend-user
Metrics:                                                 ( current / target )
  resource cpu on pods  (as a percentage of request):   <unknown> / 70%
Min replicas:                                            3
Max replicas:                                            5
Deployment pods:                                         3 current / 3 desired
Conditions:
  Type           Status  Reason              Message
  ----           ------  ------              -------
  AbleToScale    True    SucceededGetScale   the HPA controller was able to get the target's current scale
  ScalingActive  False   FailedGetResourceMetric   the HPA was unable to compute the replica count: failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch metrics from resource me
trics API: the server could not find the requested resource (get pods.metrics.k8s.io)
```



```
Events:
  Type     Reason                   Age           From                       Message
  ----     ------                   ----          ----                       -------
  Normal   SuccessfulRescale        56s           horizontal-pod-autoscaler  New size: 3; reason: Current number of replicas below Spec.MinReplicas
  Warning  FailedGetResourceMetric  11s (x3 over 41s)  horizontal-pod-autoscaler  failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch metrics from resource metrics API: t
he server could not find the requested resource (get pods.metrics.k8s.io)
  Warning  FailedComputeMetricsReplicas  11s (x3 over 41s)  horizontal-pod-autoscaler  invalid metrics (1 invalid out of 1), first error is: failed to get cpu utilization: unable to get metrics for resource cpu
: unable to fetch metrics from resource metrics API: the server could not find the requested resource (get pods.metrics.k8s.io)


Name:                                                    frontend
Namespace:                                               default
Labels:                                                  <none>
Annotations:                                             <none>
CreationTimestamp:                                       Sat, 24 Sep 2022 12:02:48 -0700
Reference:                                               Deployment/frontend
Metrics:                                                 ( current / target )
  resource cpu on pods  (as a percentage of request):   <unknown> / 70%
Min replicas:                                            3
Max replicas:                                            5
Deployment pods:                                         3 current / 3 desired
Conditions:
  Type           Status  Reason              Message
  ----           ------  ------              -------
  AbleToScale    True    SucceededGetScale   the HPA controller was able to get the target's current scale
  ScalingActive  False   FailedGetResourceMetric   the HPA was unable to compute the replica count: failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch metrics from resource me
trics API: the server could not find the requested resource (get pods.metrics.k8s.io)
Events:
  Type     Reason                   Age           From                       Message
  ----     ------                   ----          ----                       -------
  Normal   SuccessfulRescale        41s           horizontal-pod-autoscaler  New size: 3; reason: Current number of replicas below Spec.MinReplicas
  Warning  FailedGetResourceMetric  11s (x2 over 26s)  horizontal-pod-autoscaler  failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch metrics from resource metrics API: t
he server could not find the requested resource (get pods.metrics.k8s.io)
  Warning  FailedComputeMetricsReplicas  11s (x2 over 26s)  horizontal-pod-autoscaler  invalid metrics (1 invalid out of 1), first error is: failed to get cpu utilization: unable to get metrics for resource cpu
: unable to fetch metrics from resource metrics API: the server could not find the requested resource (get pods.metrics.k8s.io)


Name:                                                    reverseproxy
Namespace:                                               default
Labels:                                                  <none>
Annotations:                                             <none>
CreationTimestamp:                                       Sat, 24 Sep 2022 12:03:07 -0700
Reference:                                               Deployment/reverseproxy
Metrics:                                                 ( current / target )
  resource cpu on pods  (as a percentage of request):   <unknown> / 70%
Min replicas:                                            3
Max replicas:                                            5
Deployment pods:                                         3 current / 3 desired
Conditions:
```



```
Conditions:
  Type           Status  Reason              Message
  ----           ------  ------              -------
  AbleToScale    True    SucceededGetScale   the HPA controller was able to get the target's current scale
  ScalingActive  False   FailedGetResourceMetric   the HPA was unable to compute the replica count: failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch metrics from resource me
trics API: the server could not find the requested resource (get pods.metrics.k8s.io)
Events:
  Type     Reason                   Age           From                       Message
  ----     ------                   ----          ----                       -------
  Normal   SuccessfulRescale        41s           horizontal-pod-autoscaler  New size: 3; reason: Current number of replicas below Spec.MinReplicas
  Warning  FailedGetResourceMetric  11s (x2 over 26s)  horizontal-pod-autoscaler  failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch metrics from resource metrics API: t
he server could not find the requested resource (get pods.metrics.k8s.io)
  Warning  FailedComputeMetricsReplicas  11s (x2 over 26s)  horizontal-pod-autoscaler  invalid metrics (1 invalid out of 1), first error is: failed to get cpu utilization: unable to get metrics for resource cpu
: unable to fetch metrics from resource metrics API: the server could not find the requested resource (get pods.metrics.k8s.io)


Name:                                                    reverseproxy
Namespace:                                               default
Labels:                                                  <none>
Annotations:                                             <none>
CreationTimestamp:                                       Sat, 24 Sep 2022 12:03:07 -0700
Reference:                                               Deployment/reverseproxy
Metrics:                                                 ( current / target )
  resource cpu on pods  (as a percentage of request):   <unknown> / 70%
Min replicas:                                            3
Max replicas:                                            5
Deployment pods:                                         3 current / 3 desired
Conditions:
  Type           Status  Reason              Message
  ----           ------  ------              -------
  AbleToScale    True    SucceededGetScale   the HPA controller was able to get the target's current scale
  ScalingActive  False   FailedGetResourceMetric   the HPA was unable to compute the replica count: failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch metrics from resource me
trics API: the server could not find the requested resource (get pods.metrics.k8s.io)
Events:
  Type     Reason                   Age           From                       Message
  ----     ------                   ----          ----                       -------
  Normal   SuccessfulRescale        23s           horizontal-pod-autoscaler  New size: 3; reason: Current number of replicas below Spec.MinReplicas
  Warning  FailedGetResourceMetric  8s            horizontal-pod-autoscaler  failed to get cpu utilization: unable to get metrics for resource cpu: unable to fetch metrics from resource metrics API: the server cou
ld not find the requested resource (get pods.metrics.k8s.io)
  Warning  FailedComputeMetricsReplicas  8s       horizontal-pod-autoscaler  invalid metrics (1 invalid out of 1), first error is: failed to get cpu utilization: unable to get metrics for resource cpu: unable to f
etch metrics from resource metrics API: the server could not find the requested resource (get pods.metrics.k8s.io)
hp9ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
```

- Kubernetes `kubectl logs <your pod name>` output

```
hp@ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$ kubectl get pods
NAME                          READY   STATUS    RESTARTS       AGE
backend-feed-76c9858588-gztg6   1/1     Running   0              26s
backend-feed-76c9858588-vtr5q   1/1     Running   0              24s
backend-user-7cbd499f75-f2jst   1/1     Running   0              11s
backend-user-7cbd499f75-rwf9p   1/1     Running   0              9s
frontend-65554cd4f5-n45rj       1/1     Running   0              25h
frontend-65554cd4f5-pq7jf       1/1     Running   0              25h
reverseproxy-7c8bbd75c4-ppxg7   1/1     Running   3 (26h ago)    26h
reverseproxy-7c8bbd75c4-x97t4   1/1     Running   3 (26h ago)    26h

hp@ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$ kubectl logs -f backend-feed-76c9858588-gztg6

> udagram-api@2.0.0 prod /usr/src/app
> tsc && node ./www/server.js

Initialize database connection...
Executing (default): CREATE TABLE IF NOT EXISTS "FeedItem" ("id"   SERIAL , "caption" VARCHAR(255), "url" VARCHAR(255), "createdAt" TIMESTAMP WITH TIME ZONE, "updatedAt" TIMESTAMP WITH TIME ZONE, PRIMARY
id"));
Executing (default): SELECT i.relname AS name, ix.indisprimary AS primary, ix.indisunique AS unique, ix.indkey AS indkey, array_agg(a.attnum) as column_indexes, array_agg(a.attname) AS column_names, pg_ge
xdef(ix.indexrelid) AS definition FROM pg_class t, pg_class i, pg_index ix, pg_attribute a WHERE t.oid = ix.indrelid AND i.oid = ix.indexrelid AND a.attrelid = t.oid AND t.relkind = 'r' and t.relname = 'F
m" GROUP BY i.relname, ix.indexrelid, ix.indisprimary, ix.indisunique, ix.indkey ORDER BY i.relname;
server running http://localhost:8100
press CTRL+C to stop server
```

```
hp@ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$ kubectl get pods
NAME                          READY   STATUS    RESTARTS       AGE
backend-feed-76c9858588-5x48n   1/1     Running   0              43m
backend-feed-76c9858588-gztg6   1/1     Running   0              106m
backend-feed-76c9858588-vtr5q   1/1     Running   0              106m
backend-user-7cbd499f75-8ksd4   1/1     Running   0              43m
backend-user-7cbd499f75-f2jst   1/1     Running   0              106m
backend-user-7cbd499f75-rwf9p   1/1     Running   0              106m
frontend-65554cd4f5-j2z74       1/1     Running   0              43m
frontend-65554cd4f5-n45rj       1/1     Running   0              27h
frontend-65554cd4f5-pq7jf       1/1     Running   0              27h
reverseproxy-7c8bbd75c4-7m4gq   1/1     Running   0              42m
reverseproxy-7c8bbd75c4-ppxg7   1/1     Running   3 (28h ago)    28h
reverseproxy-7c8bbd75c4-x97t4   1/1     Running   3 (28h ago)    28h

hp@ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$ kubectl logs backend-feed-76c9858588-5x48n

> udagram-api@2.0.0 prod /usr/src/app
> tsc && node ./www/server.js

Initialize database connection...
Executing (default): CREATE TABLE IF NOT EXISTS "FeedItem" ("id"   SERIAL , "caption" VARCHAR(255), "url" VARCHAR(255), "createdAt" TIMESTAMP WITH TIME ZONE, "updatedAt" TIMESTAMP WITH TIME ZONE, PRIMARY KEY ("
id"));
Executing (default): SELECT i.relname AS name, ix.indisprimary AS primary, ix.indisunique AS unique, ix.indkey AS indkey, array_agg(a.attnum) as column_indexes, array_agg(a.attname) AS column_names, pg_get_inde
xdef(ix.indexrelid) AS definition FROM pg_class t, pg_class i, pg_index ix, pg_attribute a WHERE t.oid = ix.indrelid AND i.oid = ix.indexrelid AND a.attrelid = t.oid AND t.relkind = 'r' and t.relname = 'FeedIte
m" GROUP BY i.relname, ix.indexrelid, ix.indisprimary, ix.indisunique, ix.indkey ORDER BY i.relname;
server running http://localhost:8100
press CTRL+C to stop server

hp@ET-27608 MINGW64 ~/Desktop/project3/Microservice/deplyment (main)
$
```